

An Introduction to Machine Learning for Physicists

A Comprehensive Summary

David Hutasoit

September 9, 2026

Abstract

This article provides a comprehensive summary of *An Introduction to Machine Learning for Physicists*, a book that bridges the gap between traditional physics reasoning and modern machine learning. The book translates ML concepts into the language of statistical mechanics, information theory, and dynamical systems—entropy, energy landscapes, symmetries, and phase transitions—while providing hands-on Python implementations. Covering supervised and unsupervised learning, neural networks, generative models, probabilistic inference, reinforcement learning, and quantum machine learning, the book demonstrates that the same mathematical structures that describe physical systems also govern learning algorithms. This summary distils the key concepts, physical analogies, and practical implementations from each chapter, providing a concise reference for physicists entering the field of machine learning.

Contents

1 Introduction: The Convergence of Physics and Machine Learning

Machine learning (ML) has evolved from a specialised subfield of computer science into a transformative toolkit that permeates virtually every branch of science and technology. For physicists in particular, ML offers not just new computational tools, but a new way of thinking about data, models, and the very nature of scientific inference.

The purpose of the book is to provide an introduction to the core concepts and tools of machine learning in a manner easily understood and intuitive to physicists. While there exist many wonderful ML textbooks, they are often lengthy and use specialised language that can be unfamiliar to physicists. This book builds upon the considerable knowledge most physicists already possess in statistical physics, linear algebra, and dynamical systems in order to introduce many of the major ideas and techniques used in modern ML.

1.1 The Physicist’s Perspective

What distinguishes a physicist’s approach to ML from that of a computer scientist? The answer lies in our training. Physicists are accustomed to thinking in terms of symmetries, conservation laws, energy landscapes, and emergent phenomena. We are comfortable with coarse-graining, mean-field theories, and the renormalisation group. We understand that a model is not reality—it is an approximation, valid in some regime, that captures the essential degrees of freedom.

This perspective is remarkably well-suited to modern ML. As Mehta et al. emphasise, there are “many natural connections between ML and statistical physics”. The bias-variance tradeoff

mirrors the tradeoff between flexibility and stability in physical systems. Regularisation is the imposition of a prior—a confining potential that prevents overfitting. Gradient descent is the motion of a particle in a potential energy landscape. Neural networks are spin systems with learned couplings. Generative models are variational free energy minimisation.

1.2 Energy-Based Models and Statistical Physics

A major theme of the book is the deep connection between energy-based models in ML and statistical physics. Many ML models can be understood as Boltzmann distributions over some set of degrees of freedom:

$$p(\mathbf{x}) = \frac{1}{Z} e^{-E(\mathbf{x})}, \quad (1)$$

where $E(\mathbf{x})$ is an energy function and Z is the partition function. This formulation unifies seemingly disparate models: maximum entropy models, Restricted Boltzmann Machines, Hopfield networks, and even deep generative models can all be viewed through this lens. The goal of learning becomes the determination of an energy function such that the resulting distribution matches the data—precisely the inverse problem of statistical mechanics.

1.3 The Practical Spirit

The book emphasises hands-on, practical engagement with ML. Each chapter includes Python code examples that implement key algorithms using NumPy, scikit-learn, PyTorch, and Qiskit. The code is designed to be run and modified, encouraging experimentation and deeper understanding.

2 Chapter 1: The Intersection of Physics and Learning

This chapter establishes the conceptual groundwork for the rest of the book by translating core ML ideas into the physical language of entropy, free energy, phase transitions, and dissipative dynamics.

2.1 Information, Entropy, and Thermodynamics

At the heart of both learning and physics lies the concept of *uncertainty*. In thermodynamics, entropy quantifies the number of microscopic configurations consistent with a macroscopic state. In information theory, entropy measures the average surprise or information content of a random variable.

For a discrete random variable X , the Shannon entropy is defined as:

$$H(X) = - \sum_{x \in \mathcal{X}} p(x) \log p(x), \quad (2)$$

with the convention $0 \log 0 = 0$.

In statistical mechanics, the Gibbs entropy for a canonical ensemble is:

$$S = -k_B \sum_i p_i \ln p_i, \quad (3)$$

where $p_i = e^{-E_i/(k_B T)}/Z$ is the Boltzmann distribution. Up to the factor k_B , this is structurally identical to Shannon entropy. The maximum entropy principle—choose the distribution that maximises entropy subject to known constraints—leads naturally to the Boltzmann factor.

Jaynes showed that this principle is not merely an analogy: statistical mechanics is a form of *inference* given incomplete information. When we regularise a model, we are effectively imposing a prior distribution, which can be interpreted as adding an entropy term to the loss function.

2.2 The Loss Function as a Hamiltonian

Let $\theta \in \mathbb{R}^d$ denote the parameters of a model. The loss function $\mathcal{L}(\theta; \mathcal{D})$ plays the role of a *potential energy* or a *Hamiltonian*:

$$E(\theta) \equiv \mathcal{L}(\theta; \mathcal{D}). \quad (4)$$

The training process is a search for the global (or at least a good local) minimum of this energy landscape. This perspective allows us to import intuition from statistical mechanics—the existence of many local minima (glassy behaviour), the role of temperature (learning rate), and the possibility of phase transitions as model capacity changes.

2.3 The Learning Dynamics as Approach to Equilibrium

Gradient descent can be seen as the deterministic evolution of a particle in the potential $E(\theta)$ under high friction (overdamped Langevin dynamics):

$$\dot{\theta} = -\nabla_{\theta} E(\theta). \quad (5)$$

In the presence of stochasticity (minibatch noise), this becomes the Langevin equation:

$$d\theta_t = -\nabla_{\theta} E(\theta_t) dt + \sqrt{2T} dW_t, \quad (6)$$

where W_t is a Wiener process and T is the effective temperature set by the learning rate and batch size. The stationary distribution of this process is the Gibbs distribution $p(\theta) \propto e^{-E(\theta)/T}$.

2.4 The Bias–Variance Tradeoff

For a regression problem, the expected squared error at a test point x can be decomposed as:

$$\mathbb{E}[(y - \hat{f}(x))^2] = \underbrace{(f(x) - \mathbb{E}[\hat{f}(x)])^2}_{\text{bias}^2} + \underbrace{\mathbb{E}[(\hat{f}(x) - \mathbb{E}[\hat{f}(x)])^2]}_{\text{variance}} + \sigma^2. \quad (7)$$

Bias measures error due to overly simplistic assumptions; **variance** measures error due to excessive sensitivity to data. The physical analogy is the tradeoff between stiffness and flexibility in solid-state physics.

Practical: Estimating entropy from data using histogram binning, illustrating the bias-variance tradeoff in action.

3 Chapter 2: Supervised Learning and Energy Landscapes

Supervised learning is the task of inferring a mapping from inputs to outputs from labelled examples. This chapter treats parameters as coordinates in a high-dimensional space and the loss function as a potential energy landscape.

3.1 Linear Regression and the Quadratic Potential

With mean squared error loss, the energy function is:

$$E(\mathbf{w}, b) = \frac{1}{2N} \sum_{i=1}^N \left(y_i - (\mathbf{w}^\top \mathbf{x}_i + b) \right)^2. \quad (8)$$

This is a quadratic function—a paraboloid—with a single global minimum. The closed-form solution is:

$$\boldsymbol{\theta}^* = \left(\mathbf{X}^\top \mathbf{X} \right)^{-1} \mathbf{X}^\top \mathbf{y}. \quad (9)$$

3.2 Regularisation as a Penalty Term

Ridge (L_2) regularisation:

$$E_{\text{ridge}}(\boldsymbol{\theta}) = E(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_2^2. \quad (10)$$

The quadratic penalty creates a harmonic oscillator potential around zero, corresponding to a Gaussian prior.

Lasso (L_1) regularisation:

$$E_{\text{lasso}}(\boldsymbol{\theta}) = E(\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|_1. \quad (11)$$

The absolute value penalty produces sparse solutions, analogous to a crystal field that forces degrees of freedom to freeze out.

3.3 Gradient Descent as Dissipative Dynamics

The gradient descent update is:

$$\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k - \eta \nabla E(\boldsymbol{\theta}_k). \quad (12)$$

For a quadratic potential, convergence requires $0 < \eta < 2/\lambda_{\max}(\mathbf{H})$. Stochastic gradient descent (SGD) with minibatches introduces noise:

$$\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k - \eta \nabla E(\boldsymbol{\theta}_k) + \sqrt{2\eta T} \boldsymbol{\xi}_k, \quad (13)$$

which is the Euler–Maruyama discretisation of the Langevin equation.

3.4 Support Vector Machines and Margin Maximisation

The SVM seeks a hyperplane that maximises the margin:

$$E_{\text{SVM}}(\mathbf{w}, b) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^N \max(0, 1 - y_i(\mathbf{w}^\top \mathbf{x}_i + b)). \quad (14)$$

The kernel trick maps data into a high-dimensional feature space via $k(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$.

Practical: Linear regression with closed-form solution, gradient descent, and SGD; comparison with scikit-learn.

4 Chapter 3: Unsupervised Learning and Dimensionality

Unsupervised learning discovers hidden structure without labels. This chapter treats dimensionality reduction as finding normal modes, clustering as a first-order phase transition, and topology as a characterisation of data shape.

4.1 Principal Component Analysis (PCA)

PCA diagonalises the covariance matrix:

$$\Sigma = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i \mathbf{x}_i^\top. \quad (15)$$

The eigenvectors are the principal components; the eigenvalues represent variance captured. PCA is analogous to finding normal modes of coupled oscillators or the principal axes of inertia.

4.2 K-Means Clustering

K-means minimises the within-cluster sum of squares:

$$E = \sum_{i=1}^N \sum_{j=1}^K r_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2. \quad (16)$$

The distortion curve exhibits an elbow—a first-order phase transition—revealing the optimal number of clusters. More sophisticated criteria add a penalty term for model complexity, analogous to the free energy $F = E - TS$.

4.3 Topological Data Analysis

Persistent homology captures global shape invariants (connected components, loops, voids) that are robust to continuous deformations. This connects directly to topological phases in condensed matter physics.

Practical: PCA and K-means on the two-moons dataset; PCA for denoising.

5 Chapter 4: Neural Networks and Statistical Mechanics

This chapter explores the foundational connections between neural networks and statistical mechanics: neurons as spins, connections as exchange couplings, and dynamics as relaxation processes.

5.1 The Perceptron and the Ising Model

The perceptron is a single binary neuron:

$$y = \text{sgn}(\mathbf{w}^\top \mathbf{x} + b). \quad (17)$$

Mapping to Ising spins: $s = \text{sgn}(\tilde{\mathbf{w}}^\top \tilde{\mathbf{x}})$, with energy $E = -sh$. Learning adjusts the magnetic field.

5.2 Hopfield Networks and Spin Glasses

The Hopfield network Hamiltonian is:

$$H(\mathbf{s}) = -\frac{1}{2} \sum_{i,j} J_{ij} s_i s_j. \quad (18)$$

Hebbian learning stores patterns:

$$J_{ij} = \frac{1}{N} \sum_{\mu=1}^P \xi_i^\mu \xi_j^\mu. \quad (19)$$

The storage capacity is $\alpha_c \approx 0.138$ patterns per neuron, explained by the replica method.

5.3 Boltzmann Machines

Stochastic units updated with Glauber dynamics:

$$p(s_i = +1 | \mathbf{s}_{\setminus i}) = \frac{1}{1 + \exp\left(-2\beta \sum_j J_{ij} s_j\right)}. \quad (20)$$

The partition function is $Z = \sum_{\{\mathbf{s}\}} e^{-\beta H(\mathbf{s})}$, which is generally intractable. RBMs make learning tractable through conditional independence.

Practical: Perceptron classifier on iris data; Hopfield network for pattern completion.

6 Chapter 5: Deep Learning Architectures

This chapter scales up from shallow networks to deep architectures: MLPs, backpropagation, CNNs, and attention mechanisms.

6.1 Multi-Layer Perceptrons

MLPs are universal approximators. Each layer applies:

$$\mathbf{z}^{(\ell)} = \mathbf{W}^{(\ell)} \mathbf{h}^{(\ell-1)} + \mathbf{b}^{(\ell)}, \quad (21)$$

$$\mathbf{h}^{(\ell)} = f^{(\ell)}(\mathbf{z}^{(\ell)}). \quad (22)$$

ReLU acts like a half-harmonic potential.

6.2 Backpropagation as Adjoint Sensitivity

Backpropagation computes gradients by propagating errors backward:

$$\boldsymbol{\delta}^{(\ell)} = \left((\mathbf{W}^{(\ell+1)})^\top \boldsymbol{\delta}^{(\ell+1)} \right) \odot f'^{(\ell)}(\mathbf{z}^{(\ell)}). \quad (23)$$

This is the discrete-time analogue of adjoint sensitivity analysis.

6.3 Convolutional Neural Networks

Convolution applies a learnable kernel across spatial positions:

$$(\mathbf{X} * \mathbf{K})_{i,j} = \sum_{u,v} K_{u,v} X_{i+u,j+v}. \quad (24)$$

Translation equivariance, weight sharing, and pooling as coarse-graining mirror short-range interactions in lattice systems.

6.4 Attention Mechanisms

Scaled dot-product attention:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}} \right) \mathbf{V}. \quad (25)$$

The factor $\sqrt{d_k}$ acts as inverse temperature, and the softmax is a Boltzmann factor for all-to-all pairwise interactions.

Practical: Scaled dot-product attention; multi-head attention; MLP and CNN on MNIST; visualising CNN filters.

7 Chapter 6: Generative Models and Dynamics

Generative models learn the underlying probability distribution of data, enabling sampling and generation.

7.1 Variational Autoencoders

The ELBO is:

$$\mathcal{L} = \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}|\mathbf{z})] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z})). \quad (26)$$

This is the negative variational free energy:

$$F = \mathbb{E}_q[E] - TS(q). \quad (27)$$

The reparameterisation trick enables differentiable sampling.

7.2 Diffusion Models

Forward diffusion:

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}. \quad (28)$$

The reverse-time SDE is:

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\nabla_{\mathbf{x}} \log p(\mathbf{x}) dt + \sqrt{\beta(t)} d\mathbf{W}. \quad (29)$$

This is exactly the overdamped Langevin equation; the score $\nabla_{\mathbf{x}} \log p(\mathbf{x})$ acts as a drift.

Practical: Full DDPM implementation on 2D ring dataset; VAE implementation and latent space visualisation.

8 Chapter 7: Probabilistic Machine Learning and Bayesian Inference

This chapter makes the probabilistic framework explicit, covering Bayesian inference, Gaussian processes, and variational inference.

8.1 Bayesian Inference

Bayes' theorem:

$$p(\boldsymbol{\theta}|\mathcal{D}) = \frac{p(\mathcal{D}|\boldsymbol{\theta}) p(\boldsymbol{\theta})}{p(\mathcal{D})}. \quad (30)$$

The posterior is structurally identical to the Boltzmann distribution; the marginal likelihood is the partition function.

8.2 Bayesian Linear Regression

The posterior mean and covariance are:

$$\boldsymbol{\mu}_N = \mathbf{S}_N \boldsymbol{\Phi}^\top \mathbf{y} / \sigma^2, \quad (31)$$

$$\mathbf{S}_N^{-1} = \mathbf{S}_0^{-1} + \boldsymbol{\Phi}^\top \boldsymbol{\Phi} / \sigma^2. \quad (32)$$

The predictive distribution has variance $\sigma^2 + \boldsymbol{\phi}_*^\top \mathbf{S}_N \boldsymbol{\phi}_*$.

8.3 Gaussian Processes

GPs place a prior directly on functions:

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}')). \quad (33)$$

The RBF kernel is:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2}\right). \quad (34)$$

The predictive mean and covariance are computed from the joint Gaussian distribution.

8.4 Variational Inference

VI approximates the posterior by minimising the KL divergence:

$$q^*(\boldsymbol{\theta}) = \arg \min_{q \in \mathcal{Q}} D_{\text{KL}}(q(\boldsymbol{\theta}) \parallel p(\boldsymbol{\theta}|\mathcal{D})). \quad (35)$$

This is equivalent to maximising the ELBO.

Practical: Gaussian process regression with scikit-learn; Bayesian linear regression with polynomial features.

9 Chapter 8: Reinforcement Learning and Optimal Control

Reinforcement learning formulates sequential decision-making as optimal control under uncertainty.

9.1 Markov Decision Processes

An MDP is defined by $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$. The goal is to maximise the expected discounted return:

$$G_t = \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \right]. \quad (36)$$

9.2 The Bellman Equation

The Bellman equation is:

$$V^\pi(s) = \sum_a \pi(a|s) \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V^\pi(s')]. \quad (37)$$

This is the discrete-time Hamilton-Jacobi-Bellman equation.

9.3 Q-Learning and DQN

Q-learning updates:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]. \quad (38)$$

DQN uses experience replay and target networks for stability.

9.4 Exploration-Exploitation

Boltzmann exploration:

$$\pi(a|s) = \frac{e^{Q(s,a)/T}}{\sum_b e^{Q(s,b)/T}}. \quad (39)$$

Annealing T over time mimics simulated annealing.

Practical: DQN on CartPole using PyTorch.

10 Chapter 9: Quantum Machine Learning

QML leverages the exponential growth of Hilbert space to represent and process information.

10.1 Qubits and Entanglement

A qubit state is:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle, \quad |\alpha|^2 + |\beta|^2 = 1. \quad (40)$$

For n qubits, the Hilbert space dimension is 2^n . Entanglement is the essential resource.

10.2 Parameterised Quantum Circuits

A PQC defines:

$$f(\boldsymbol{\theta}; \mathbf{x}) = \langle 0|U^\dagger(\boldsymbol{\theta})M(\mathbf{x})U(\boldsymbol{\theta})|0\rangle. \quad (41)$$

The parameter shift rule computes gradients:

$$\frac{\partial}{\partial \theta_l} \langle O \rangle = \frac{1}{2} (\langle O \rangle_{\theta_l + \pi/2} - \langle O \rangle_{\theta_l - \pi/2}). \quad (42)$$

10.3 Quantum Kernels

Quantum kernels compute overlaps in Hilbert space:

$$k(\mathbf{x}, \mathbf{x}') = |\langle \phi(\mathbf{x}) | \phi(\mathbf{x}') \rangle|^2. \quad (43)$$

Practical: Variational quantum classifier on concentric circles using Qiskit.

11 Chapter 10: Case Studies in Physics

11.1 Particle Physics: LHC Event Classification

Signal (Higgs) distinguished from background (QCD jets) using deep learning. ML achieves AUC > 0.95 , outperforming traditional methods.

11.2 Astrophysics: Gravitational Wave Detection

CNNs detect signals buried in non-Gaussian noise in real time with $>95\%$ accuracy, offering orders-of-magnitude speedup over matched filtering.

11.3 Condensed Matter: Phase Transitions

PCA on Ising spin configurations reveals the order parameter and critical temperature without prior knowledge.

11.4 Quantum Information: State Tomography

Neural networks reconstruct quantum states from limited measurements, requiring far fewer resources than traditional tomography.

Practical: Code examples for each case study.

12 Chapter 11: Challenges and Future Directions

12.1 Key Challenges

- **Interpretability:** LIME, SHAP, saliency maps, and attention visualisation
- **Data Scarcity:** Transfer learning, few-shot learning, generative models for augmentation
- **Physics-Informed ML:** PINNs embed physical laws directly into the loss function
- **Hardware:** GPU/TPU scaling, quantum hardware for specific tasks

12.2 Future Directions

- Foundation Models for Physics (pre-trained on simulations)
- Automated Scientific Discovery (symbolic regression)
- Self-Supervised Learning on unlabelled data
- Digital Twins for real-time simulation and optimisation
- Interdisciplinary collaboration across physics, CS, and AI

Practical: LIME for MNIST; PINN for Burgers equation; few-shot learning with prototypical networks.

13 Conclusion: The Road Ahead

The book establishes that physics provides a natural language for machine learning. Entropy and free energy unify uncertainty; energy landscapes describe optimisation; symmetries provide inductive biases; and statistical mechanics explains learning dynamics.

13.1 Key Lessons

1. **Physics Provides a Natural Language for ML** – The conceptual toolkit of physics maps beautifully onto ML concepts.
2. **Uncertainty Matters** – Bayesian methods, Gaussian processes, and variational inference quantify uncertainty.
3. **Symmetry Is a Powerful Inductive Bias** – CNNs, equivariant networks, and gauge invariance impose physical symmetries.
4. **Models Are Not Reality** – Every model is an approximation; validation and understanding limitations are essential.
5. **The Best Models Combine Physics and Data** – Physics-informed approaches achieve more with less data.
6. **The Future Is Interdisciplinary** – The most exciting advances will come at the intersections.

13.2 Final Reflections

Machine learning is not a replacement for physics—it is an augmentation. The great discoveries of the future will come from the synergy between human intuition and computational power. As a physicist, you bring a deep understanding of symmetries, conservation laws, and physical

principles—knowledge that is invaluable in building models that are not just accurate, but also interpretable and physically meaningful.

We stand at the threshold of a new era, where ML will help us probe the nature of reality, from the subatomic to the cosmic scale. The journey is just beginning.

“The important thing is not to stop questioning. Curiosity has its own reason for existing.”
— Albert Einstein

Core Physical Analogies

Machine Learning Concept	Physical Analogy
Loss function	Hamiltonian / potential energy
Gradient descent	Particle in a viscous medium
SGD noise	Thermal fluctuations (Langevin dynamics)
Regularisation	Confining potential (prior)
Neural networks	Spin systems (Ising model)
Backpropagation	Adjoint sensitivity analysis
CNNs	Short-range interactions / lattices
Attention	All-to-all interactions / Boltzmann factor
VAEs	Variational free energy minimisation
Diffusion models	Langevin dynamics / cooling
Bayesian inference	Boltzmann distribution
RL exploration	Simulated annealing
PCA	Normal modes / principal axes
K-means	First-order phase transition

Bibliography

References

- [1] Mehta, P., Bukov, M., Wang, C.-H., Day, A. G. R., Richardson, C., Fisher, C. K., & Schwab, D. J. (2019). *A high-bias, low-variance introduction to Machine Learning for physicists*. Physics Reports, 810, 1-124.
- [2] Hopfield, J. J. (1982). *Neural networks and physical systems with emergent collective computational abilities*. Proceedings of the National Academy of Sciences, 79(8), 2554-2558.
- [3] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- [4] Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer.
- [5] Murphy, K. P. (2012). *Machine Learning: A Probabilistic Perspective*. MIT Press.
- [6] Carleo, G., Cirac, I., Cranmer, K., Daudet, L., Schuld, M., Tishby, N., Vogt-Maranto, L., & Zdeborová, L. (2019). *Machine learning and the physical sciences*. Reviews of Modern Physics, 91(4), 045002.
- [7] Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). *Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations*. Journal of Computational Physics, 378, 686-707.
- [8] Karniadakis, G. E., Kevrekidis, I. G., Lu, L., Perdikaris, P., Wang, S., & Yang, L. (2021). *Physics-informed machine learning*. Nature Reviews Physics, 3(6), 422-440.